

## DAY- 4

# SQL SUBQUERIES & SET OPERATIONS - Study & Revision Notes

### ◇ Introduction

#### What is a Subquery?

A **subquery** is a **query inside another query**. It is executed first, and its result is used by the outer query.

#### Types of Subqueries

1. **Single-row subquery** → Returns **one value**.
2. **Multiple-row subquery** → Returns **multiple values**.
3. **Correlated subquery** → Inner query depends on the outer query.
4. **Nested subquery** → Subquery inside another subquery.
5. **Set operations** → Combines results from multiple queries.

## 1. Single-Row Subqueries

Returns **one value** and is used with comparison operators like =, <, >, etc.

### ☒ Find the Employee with the Lowest Salary

```
sql
CopyEdit
SELECT ENAME, SAL
FROM SCOTT.EMP
WHERE SAL = (SELECT MIN(SAL) FROM SCOTT.EMP);
```

◆ **Explanation:**

- (SELECT MIN(SAL) FROM SCOTT.EMP) → Finds **lowest salary**.
- SAL = (subquery result) → Matches employee **with that salary**.

☑ **Find Employees Earning More than the Average Salary**

```
sql
CopyEdit
SELECT ENAME, SAL
FROM SCOTT.EMP
WHERE SAL > (SELECT AVG(SAL) FROM SCOTT.EMP);
```

◆ **Explanation:**

- (SELECT AVG(SAL) FROM SCOTT.EMP) → Finds **average salary**.
- SAL > (subquery result) → Filters **employees earning above average**.

## ✂ 2. Multiple-Row Subqueries

Returns **multiple values** and uses operators like IN, ANY, ALL.

✗ **Incorrect Query (Error: Subquery returns multiple rows)**

```
sql
CopyEdit
SELECT *
FROM SCOTT.EMP
WHERE SAL > (SELECT SAL FROM SCOTT.EMP WHERE ENAME='ALLEN' OR
ENAME='JONES');
```

◆ **Problem:**

- The subquery returns **multiple salaries**.
- The > operator **expects only one value**.

## ☒ Correct Query Using ALL

```
sql
CopyEdit
SELECT *
FROM SCOTT.EMP
WHERE SAL > ALL (SELECT SAL FROM SCOTT.EMP WHERE ENAME='ALLEN' OR
ENAME='JONES');
```

### ◆ Explanation:

- SAL > ALL(subquery) → Finds employees **earning more than both Allen and Jones**.

## ☒ Find Employees Working in King's or Scott's Department

```
sql
CopyEdit
SELECT *
FROM SCOTT.EMP
WHERE DEPTNO = ANY (SELECT DEPTNO FROM SCOTT.EMP WHERE ENAME IN
('KING', 'SCOTT'));
```

### ◆ Explanation:

- (SELECT DEPTNO FROM SCOTT.EMP WHERE ENAME IN ('KING', 'SCOTT')) → Finds **departments of King & Scott**.
- DEPTNO = ANY (subquery result) → Filters **employees working in those departments**.

### ◆ Alternate Query Using IN (More Readable)

```
sql
CopyEdit
SELECT *
FROM SCOTT.EMP
WHERE DEPTNO IN (SELECT DEPTNO FROM SCOTT.EMP WHERE ENAME IN
('KING', 'SCOTT'));
```

## ✂ 3. Using Subqueries for Additional Columns

```
sql
CopyEdit
SELECT SCOTT.EMP.*,
       (SELECT MIN(SAL) FROM SCOTT.EMP) AS MIN_SAL,
       (SELECT MAX(SAL) FROM SCOTT.EMP) AS MAX_SAL
FROM SCOTT.EMP;
```

### ◆ Explanation:

- Adds **MIN\_SAL** and **MAX\_SAL** as additional columns.
- The subqueries return **single values**, so they can be used in SELECT.

## ✂ 4. Using Subqueries in FROM Clause

### ☒ Find Salary Statistics per Department

```
sql
CopyEdit
SELECT *
FROM (SELECT DEPTNO,
             MIN(SAL) AS MIN_SAL,
             MAX(SAL) AS MAX_SAL,
             ROUND(AVG(SAL),2) AS AVG_SAL,
             SUM(SAL) AS TOTAL_SAL
      FROM SCOTT.EMP
      GROUP BY DEPTNO);
```

### ◆ Explanation:

- Groups by DEPTNO and calculates **salary statistics**.
- This inner query is treated as a **temporary table**.

## ✂ 5. WITH Clause (Common Table Expressions - CTE)

sql

CopyEdit

```
WITH EMP_SUMMARY AS (  
    SELECT DEPTNO,  
           MIN(SAL) AS MIN_SAL,  
           MAX(SAL) AS MAX_SAL,  
           ROUND(AVG(SAL),2) AS AVG_SAL,  
           SUM(SAL) AS TOTAL_SAL  
    FROM SCOTT.EMP  
    GROUP BY DEPTNO)  
SELECT * FROM SCOTT.EMP JOIN EMP_SUMMARY USING(DEPTNO) ORDER BY  
DEPTNO;
```

### ◆ Explanation:

- WITH EMP\_SUMMARY AS (...) creates a **temporary result set**.
- Joins the EMP table with EMP\_SUMMARY on DEPTNO.

## ✂ 6. EXISTS vs NOT EXISTS

### ☒ Find Departments with Employees

sql

CopyEdit

```
SELECT *  
FROM SCOTT.DEPT  
WHERE EXISTS (SELECT * FROM SCOTT.EMP WHERE SCOTT.EMP.DEPTNO =  
SCOTT.DEPT.DEPTNO);
```

### ◆ Explanation:

- If employees exist in a department, it is **included** in the result.

### ☒ Find Departments Without Employees

sql  
CopyEdit  
SELECT \*  
FROM SCOTT.DEPT  
WHERE NOT EXISTS (SELECT \* FROM SCOTT.EMP WHERE SCOTT.EMP.DEPTNO =  
SCOTT.DEPT.DEPTNO);

#### ◆ Explanation:

- If no employees exist in a department, it is **included**.

## ✂ 7. Correlated Subqueries

A **correlated subquery** depends on the outer query.

### ☒ Find Employees Earning More than the Average Salary of Their Department

sql  
CopyEdit  
SELECT \*  
FROM SCOTT.EMP e1  
WHERE SAL > (SELECT AVG(e2.SAL) FROM SCOTT.EMP e2 WHERE e1.DEPTNO =  
e2.DEPTNO);

#### ◆ Explanation:

- The inner query calculates **average salary for the same department**.
- Each employee's salary is compared **within their department**.

### ☒ Find the Highest-Paid Employee in Each Department

sql  
CopyEdit

```
SELECT *
FROM SCOTT.EMP e1
WHERE e1.SAL = (SELECT MAX(e2.SAL) FROM SCOTT.EMP e2 WHERE e1.DEPTNO
= e2.DEPTNO);
```

◆ **Explanation:**

- The subquery finds **max salary per department**.
- The outer query filters employees **who match the max salary**.

## ✂ 8. SET OPERATIONS

### (Combining Results from Multiple Queries)

#### ☒ Creating Sample Tables

```
sql
CopyEdit
CREATE TABLE emp1(empno NUMBER, ename VARCHAR2(20), deptno NUMBER);
INSERT INTO emp1 VALUES(1001, 'Ajit',10);
INSERT INTO emp1 VALUES(1002, 'Bob',10);
INSERT INTO emp1 VALUES(1003, 'John',20);
INSERT INTO emp1 VALUES(1004, 'Els',20);

CREATE TABLE emp2(empno NUMBER, ename VARCHAR2(20), deptno NUMBER);
INSERT INTO emp2 VALUES(1001, 'John',20);
INSERT INTO emp2 VALUES(1002, 'Elsa',20);
INSERT INTO emp2 VALUES(1003, 'Raju',30);
INSERT INTO emp2 VALUES(1004, 'Sunny',30);
```

## ✂ 9. UNION (Removes Duplicates)

```
sql
CopyEdit
SELECT * FROM emp1
UNION
```

```
SELECT * FROM emp2;
```

- ◆ **Combines both tables, removes duplicates.**

## ✂ 10. UNION ALL (Includes Duplicates)

```
sql
CopyEdit
SELECT * FROM emp1
UNION ALL
SELECT * FROM emp2;
```

- ◆ **Combines both tables, keeps duplicates.**

## ✂ 11. INTERSECT (Find Common Records)

```
sql
CopyEdit
SELECT * FROM emp1
INTERSECT
SELECT * FROM emp2;
```

- ◆ **Returns only matching rows** from both tables.

## ✂ 12. MINUS (Find Unique Records in emp1)

```
sql
CopyEdit
SELECT * FROM emp1
MINUS
SELECT * FROM emp2;
```



- ◆ Returns records in emp1 that are NOT in emp2.

## Summary

- **Subqueries** extract data dynamically.
- **Single-row vs Multiple-row subqueries.**
- **Correlated subqueries** iterate per row.
- **EXISTS vs NOT EXISTS** filters presence/absence.
- **Set operations (UNION, INTERSECT, MINUS)** combine results.

## SQL HIERARCHICAL RETRIEVAL - Study & Revision Notes

### ◇ Introduction

**Hierarchical queries** are used to retrieve data that is stored in a tree-like structure. In **Oracle SQL**, the **CONNECT BY** clause is used to define parent-child relationships.

### Key Concepts

- **CONNECT BY PRIOR column = column** → Defines the parent-child relationship.
- **START WITH condition** → Defines the root node.
- **LEVEL** → Represents the depth in the hierarchy.
- **ORDER SIBLINGS BY** → Sorts records within the same level.
- **SYS\_CONNECT\_BY\_PATH** → Shows the path from root to each node.
- **CONNECT\_BY\_ISLEAF** → Identifies leaf nodes (employees with no subordinates).

## ✂ 1. Basic Hierarchical Query

### ☑ Find the Hierarchy Under 'KING'

```
sql
CopyEdit
SELECT EMPNO, ENAME, MGR
FROM SCOTT.EMP
CONNECT BY PRIOR EMPNO = MGR
START WITH ENAME = 'KING';
```

#### ◆ Explanation:

- **START WITH ENAME='KING'** → Starts hierarchy from King.
- **CONNECT BY PRIOR EMPNO = MGR** → Builds hierarchy by linking MGR (manager ID) to EMPNO (employee ID).
- Returns **all employees reporting directly or indirectly** to King.

## ✂ 2. Adding Levels in Hierarchy

```
sql
CopyEdit
SELECT EMPNO, ENAME, MGR, LEVEL
FROM SCOTT.EMP
CONNECT BY PRIOR EMPNO = MGR
START WITH ENAME = 'KING'
ORDER BY LEVEL;
```

#### ◆ Explanation:

- **LEVEL** → Represents the depth of the employee in the hierarchy.
- **ORDER BY LEVEL** → Sorts the output based on levels.

### ☑ Output Example:

EMPNO	ENAME	MGR	LEVEL
7839	KING	NULL	1
7566	JONES	7839	2

7698	BLAKE	7839	2
7782	CLARK	7839	2

### ✂ 3. Sorting Siblings Alphabetically

sql

CopyEdit

```
SELECT EMPNO, ENAME, MGR, LEVEL
FROM SCOTT.EMP
CONNECT BY PRIOR EMPNO = MGR
START WITH ENAME = 'KING'
ORDER SIBLINGS BY ENAME;
```

#### ◆ Explanation:

- ORDER SIBLINGS BY ENAME → Sorts **employees at the same hierarchy level** alphabetically.

#### ✓ Example Output (Sorted Alphabetically by Name)

EMPNO	ENAME	MGR	LEVEL
7839	KING	NULL	1
7698	BLAKE	7839	2
7782	CLARK	7839	2
7566	JONES	7839	2

### ✂ 4. Sorting Siblings in Descending Order

sql

CopyEdit

```
SELECT EMPNO, ENAME, MGR, LEVEL
FROM SCOTT.EMP
CONNECT BY PRIOR EMPNO = MGR
START WITH ENAME = 'KING'
ORDER SIBLINGS BY ENAME DESC;
```

#### ◆ Explanation:

- ORDER SIBLINGS BY ENAME DESC → Sorts **siblings in reverse alphabetical order**.

✅ **Example Output (Sorted in Reverse Order)**

EMPNO	ENAME	MGR	LEVEL
7839	KING	NULL	1
7566	JONES	7839	2
7782	CLARK	7839	2
7698	BLAKE	7839	2

## 🔧 5. Displaying the Path in Hierarchy

```
sql
CopyEdit
SELECT SYS_CONNECT_BY_PATH(ENAME, '/')
FROM SCOTT.EMP
CONNECT BY PRIOR EMPNO = MGR
START WITH ENAME = 'KING';
```

◆ **Explanation:**

- **SYS\_CONNECT\_BY\_PATH(column, delimiter)** → Displays **full path** from root to each employee.
- **Example Output:**

```
swift
CopyEdit
/KING
/KING/JONES
/KING/JONES/SCOTT
/KING/BLAKE
/KING/CLARK
```

📌 **Use Case:** Helps visualize reporting structure.

## 🔧 6. Identifying Leaf Nodes

sql

CopyEdit

```
SELECT ENAME, CONNECT_BY_ISLEAF
FROM SCOTT.EMP
CONNECT BY PRIOR EMPNO = MGR
START WITH ENAME = 'KING';
```

### ◆ Explanation:

- `CONNECT_BY_ISLEAF` → Returns 1 if the employee has **no subordinates**, otherwise 0.

### ✅ Example Output:

ENAME	CONNECT_BY_ISLEAF
KING	0
JONES	0
SCOTT	1
BLAKE	0
CLARK	1

🔧 **Use Case:** Helps find employees who are **not managers**.

## 🚀 Summary

- **Hierarchical queries** are useful for **organizational structures**.
- `START WITH` → Defines the **root** of the hierarchy.
- `CONNECT BY PRIOR` → Establishes **parent-child relationships**.
- `LEVEL` → Represents **depth** in the hierarchy.
- `ORDER SIBLINGS BY` → Sorts **employees at the same level**.
- `SYS_CONNECT_BY_PATH` → Displays **full hierarchy path**.
- `CONNECT_BY_ISLEAF` → Identifies **employees with no subordinates**.

# SQL Views - Study & Revision Notes

## ◇ Introduction to Views

- **What is a View?**
  - A **view** is a **virtual table** based on the result set of an SQL query.
  - It does **not store data** itself but displays data stored in other tables.
- **Why Use Views?**
  - **Fine-grained access control:** Restrict access to specific data.
  - **Simplify complex queries:** Abstract complicated joins or calculations.
  - **Data security:** Hide sensitive data.
- **Types of Views**
  - **Simple View:** Based on **one table**, no functions, aggregations, or joins.
  - **Complex View:** Based on **multiple tables** or uses functions, aggregations, or joins.

## ✂ 1. Creating and Inserting into a Table

### Create EMP Table

```
sql
CopyEdit
CREATE TABLE EMP(
    EMPNO NUMBER(4,0) NOT NULL,
    ENAME VARCHAR2(10),
    JOB VARCHAR2(9),
    MGR NUMBER(4,0),
    HIREDATE DATE,
    SAL NUMBER(7,2),
    COMM NUMBER(7,2),
    DEPTNO NUMBER(2,0)
);
```

#### ◆ Explanation:

- **Defines the EMP table** with columns for employee details.

## Insert Data from SCOTT.EMP

sql

CopyEdit

```
INSERT INTO EMP (SELECT * FROM SCOTT.EMP);
```

### ◆ Explanation:

- **Copies data** from the SCOTT . EMP table into the new EMP table.

## ✂ 2. Creating a Simple View

### Create a View for Department 10

sql

CopyEdit

```
CREATE VIEW DEPT_10_VIEW AS SELECT * FROM EMP WHERE DEPTNO=10;
```

### ◆ Explanation:

- **Simple view** showing all employees in department 10.

### Modify the View

sql

CopyEdit

```
CREATE OR REPLACE VIEW DEPT_10_VIEW AS SELECT * FROM EMP WHERE DEPTNO=10;
```

### ◆ Explanation:

- **Replaces** the existing view if it exists, or **creates** it if not.

## ✂ Attempt to Alter a View

sql

CopyEdit

```
ALTER VIEW DEPT_10_VIEW ADD TEST_COL NUMBER;
```

#### ✖ Error:

- **Views cannot be altered** to add columns. You must recreate them.

## ✂ 3. Performing DML Operations on Views

### Insert into the View

sql

CopyEdit

```
INSERT INTO DEPT_10_VIEW VALUES(9999, 'DNYANESH', 'ANALYST', 7839, '04-JUL-22', 1000, NULL, 10);
```

#### ♦ Explanation:

- Inserts a **new employee** into DEPT\_10\_VIEW.
- **Changes reflect** in the base EMP table since it's a **simple view**.

### ✂ Inserting with Incorrect Department

sql

CopyEdit

```
INSERT INTO DEPT_10_VIEW VALUES(8888, 'RAHUL', 'ANALYST', 7839, '04-JUL-22', 1000, NULL, 20);
```

#### ✖ Error:

- Department 20 **violates the view condition** (DEPTNO=10).

### Delete from the View

sql

CopyEdit

```
DELETE FROM DEPT_10_VIEW WHERE EMPNO=8888;
```

#### ♦ Explanation:



- **Deletes the record** both from the view and the base table.

## ✂ 4. Restricting DML with View Options

### Read-Only View

sql

CopyEdit

```
CREATE OR REPLACE VIEW DEPT_10_VIEW AS SELECT * FROM EMP WHERE  
DEPTNO=10 WITH READ ONLY;
```

#### ◆ Explanation:

- **Prevents any DML** operations (INSERT, UPDATE, DELETE).

### ✂ Attempting to Insert into Read-Only View

sql

CopyEdit

```
INSERT INTO DEPT_10_VIEW VALUES(9999, 'DNYANESH', 'ANALYST', 7839,  
'04-JUL-22', 1000, NULL, 10);
```

#### ✂ Error:

- **Cannot perform DML** on a read-only view.

### Using WITH CHECK OPTION

sql

CopyEdit

```
CREATE OR REPLACE VIEW DEPT_10_VIEW AS SELECT * FROM EMP WHERE  
DEPTNO=10 WITH CHECK OPTION;
```

#### ◆ Explanation:

- Ensures that **all DML operations** conform to the **view's WHERE clause**.

## ✖ Inserting a Record that Violates the View Condition

sql

CopyEdit

```
INSERT INTO DEPT_10_VIEW VALUES(7777, 'RAHUL', 'ANALYST', 7839, '04-JUL-22', 1000, NULL, 20);
```

### ✖ Error:

- **Rejected** because DEPTNO=20 does **not match** the view's condition (DEPTNO=10).

## ✖ 5. Updating Data via Views

### Updating Salary in the Base Table

sql

CopyEdit

```
UPDATE EMP SET SAL = SAL * 1.10 WHERE ENAME = 'KING';
```

- **Increases KING's salary by 10%.**

### Verify Update in Both View and Table

sql

CopyEdit

```
SELECT * FROM DEPT_10_VIEW WHERE ENAME = 'KING';  
SELECT * FROM EMP WHERE ENAME = 'KING';
```

- **Check both the view and the table** to confirm the update.

## ✖ 6. Creating a Complex View

### Create an Aggregate View

sql

CopyEdit

```
CREATE OR REPLACE VIEW EMP_SUMMARY AS
SELECT DEPTNO,
       COUNT(*) AS TOTAL_EMPLOYEES,
       MIN(SAL) AS MIN_SAL,
       MAX(SAL) AS MAX_SAL,
       SUM(SAL) AS TOTAL_SAL,
       AVG(SAL) AS AVG_SAL
FROM EMP
GROUP BY DEPTNO;
```

◆ **Explanation:**

- **Aggregates data** by department.
- **Complex view** due to aggregation.

## ✗ Attempt to Update a Complex View

```
sql
CopyEdit
UPDATE EMP_SUMMARY SET TOTAL_EMPLOYEES = 10 WHERE DEPTNO = 30;
```

✗ **Error:**

- **DML operations** are **not allowed** on views with aggregations.

## 🔗 7. Combining Data from Multiple Tables

### Join View

```
sql
CopyEdit
CREATE VIEW EMP_DEPT_DATA AS
SELECT * FROM SCOTT.EMP JOIN SCOTT.DEPT USING(DEPTNO);
```

◆ **Explanation:**

- **Combines employee and department data** in one view.
- **Useful** for seeing related data together.

## Summary

- **Views** are flexible tools for data abstraction and security.
- **Simple Views:** Allow basic DML operations.
- **Complex Views:** Typically **read-only** due to aggregations or joins.
- **WITH CHECK OPTION:** Ensures data integrity in DML operations.
- **Read-Only Views:** Prevent any data modifications.

## SQL Query Optimization & EXPLAIN PLAN - Study & Revision Notes

This guide covers **query execution plans**, **indexes**, and **performance optimization** using EXPLAIN PLAN in Oracle SQL.

### ◇ Understanding EXPLAIN PLAN

- EXPLAIN PLAN **shows how Oracle executes a query**.
- Helps in **query tuning** by identifying **full table scans**, **index usage**, and **joins**.
- **DBMS\_XPLAN.DISPLAY()** retrieves the **execution plan** from memory.

## 1. Basic Query Execution

sql

CopyEdit

```
SELECT * FROM EMP;
```

- ◆ **Retrieves all rows** from the EMP table.

## Check Execution Plan

sql

CopyEdit

```
EXPLAIN PLAN FOR SELECT * FROM EMP WHERE DEPTNO=10;  
SELECT * FROM DBMS_XPLAN.DISPLAY();
```

### ◆ Breakdown:

1. **EXPLAIN PLAN FOR** - Prepares the query execution plan.
2. **DBMS\_XPLAN.DISPLAY()** - Displays the execution plan.

### 📌 Key Observations:

- If there is **no index on DEPTNO**, the query may **perform a full table scan**.
- A **full table scan is slow** when the table has many rows.

## 🔧 2. Checking Synonyms

sql

CopyEdit

```
SELECT * FROM USER_SYNONYMS;
```

- ◆ **Lists all synonyms** (aliases) in the user's schema.

## 🔧 3. Creating and Dropping an Index

sql

CopyEdit

```
CREATE INDEX DEPT_IDX ON EMP(DEPTNO);
```

- ◆ **Creates an index** on the DEPTNO column.

sql

CopyEdit

```
DROP INDEX DEPT_IDX;
```

- ◆ Removes the index.

#### ✚ Why Create an Index?

- **Speeds up queries** that use WHERE DEPTNO=10 by **avoiding full table scans**.
- **Indexes work best** when filtering a small subset of data.

## ✚ 4. Checking Query Optimization with Index

sql

CopyEdit

```
EXPLAIN PLAN FOR SELECT * FROM EMP WHERE DEPTNO=10;  
SELECT * FROM DBMS_XPLAN.DISPLAY();
```

- ◆ After creating an index, this query should use an **INDEX SCAN** instead of a **FULL TABLE SCAN**.

## ✚ 5. Checking Execution Plan for JOB Column

sql

CopyEdit

```
EXPLAIN PLAN FOR SELECT * FROM EMP WHERE JOB='MANAGER';  
SELECT * FROM DBMS_XPLAN.DISPLAY();
```

#### ✚ Key Notes:

- If JOB is **not indexed**, it may still perform a **full table scan**.
- **Indexing JOB could improve performance** if filtering on this column is frequent.

## ✚ 6. Checking Execution Plan for CROSS JOIN

sql

CopyEdit

```
EXPLAIN PLAN FOR SELECT * FROM SCOTT.EMP CROSS JOIN SCOTT.DEPT;  
SELECT * FROM DBMS_XPLAN.DISPLAY();
```

#### Key Observations:

- **CROSS JOIN produces a Cartesian product** (every row in EMP is combined with every row in DEPT).
- **Inefficient for large tables.** Use **INNER JOIN** or filters to optimize.

## 7. Checking Execution Plan for Aggregations

sql

CopyEdit

```
EXPLAIN PLAN FOR  
SELECT DEPTNO, COUNT(*) AS TOTAL_EMPLOYEES,  
       MIN(SAL) AS MIN_SAL,  
       MAX(SAL) AS MAX_SAL,  
       SUM(SAL) AS TOTAL_SAL,  
       AVG(SAL) AS AVG_SAL  
FROM EMP  
GROUP BY DEPTNO;  
SELECT * FROM DBMS_XPLAN.DISPLAY();
```

#### Key Observations:

- **Uses GROUP BY**, which may trigger a **SORT operation**.
- **Indexes may not help much** in aggregation queries.
- **Adding indexes on frequently grouped columns** (DEPTNO) may improve performance.

## Summary

- **EXPLAIN PLAN** helps analyze query performance.
- **Indexes** improve filtering but are ineffective for aggregation.
- **Avoid CROSS JOINS** unless necessary.

- **Always check execution plans** before optimizing queries.